



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020




Packet mark in a Cloud Native world

Joe Stringer

[Cilium.io](https://cilium.io)



the internet is held together with 

the internet is held together with duct tape	Remove
the internet is held together with bubble gum	Remove
the internet is held together with baling wire	Remove
the internet is held together with popsicle sticks	Remove
the internet is held together with pixie dust	Remove
the internet is held together with prayers	Remove
the internet is held together with skb->mark	Remove

Overview

1 Background

2 Use cases

3 Observations & Challenges

Mark of the

- fw_mark
- skb_mark
- ct_mark
- SO_MARK
- xfrm_mark
- pkt_mark

```
struct sk_buff {  
    ...  
    __u32 mark;  
    ...  
}
```

So what does the mark represent?

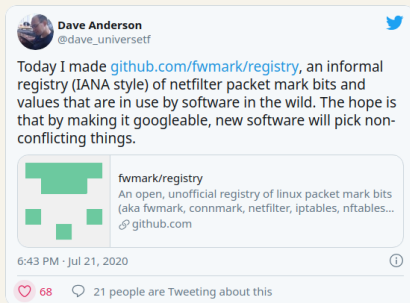
- Nothing..

So what does the mark represent?

- Nothing..
- Anything!

So what does the mark represent?

- Nothing..
- Anything!
- **MAGIC.** ✨



https://twitter.com/dave_universetf/status/1285752332135788544

eBPF-based Networking, Observability, and Security

 Follow @ciliumproject  Join Our Slack



Networking Observability Security



Highly Scalable
Kubernetes CNI



Identity-aware Network
Visibility



Advanced Network
Policy



Kube-proxy Load
Balancer Replacement



Network Metrics +
Troubleshooting



Security Forensics +
Audit



Multi-cluster
Connectivity



API-aware Network
Observability



Transparent Encryption

Cloud Native networking



ANTREA



Methodology

- 1 Look at CNCF landscape¹
- 2 Find the project on GitHub
- 3 Search for \$mark_name
- 4 ???
- 5 Knowledge!

¹<https://landscape.cncf.io/category=cloud-native-network&format=card-mode&grouping=category>

Use cases

Network policy

- 1 bit, two variations:
 - 1 bit -> drop²
 - 1 bit -> allow
- Store complex path through rules into mark
- Typically netfilter -> netfilter



² Kubernetes default

Transparent encryption

- 2+ bits
 - 1 bit encrypt, 1 bit decrypt
 - Variation: key selector
- { eBPF, netfilter } -> xfrm



Virtual IP services

- 1+ bits, request DNAT
 - 1 bit: route towards bridge for DNAT
 - 30 bits representing hashed 3-tuple
- { eBPF, netfilter } -> routing -> netfilter
- OVS -> routing -> OVS



IP masquerade

- 1+ bits, request SNAT
 - Variation: 1 bit, Skip SNAT
 - Variation: 32 bits for source address selection
- Connection may not originate on the node
- {eBPF, OVS, netfilter} -> netfilter



Multi-homing

- 1 bit, two variations:
 - Reply via primary device
 - Default: Pod communicates via secondary device
 - Inbound connections must reply via primary device
 - Store & restore in connmark
 - Route via management interface
- { socket, netfilter } -> routing



Application identity

- Variable bits
 - 4 bit pattern: “local” traffic
 - 16+ bits: Carry Identity to destination
 - Policy routing
 - Portmap plugin
- { eBPF, netfilter } -> routing -> eBPF



Service proxy

- 1+ bits depending on context
 - 1 bit, route locally
 - 16 bit tproxy port towards proxy
 - 16+ bit Identity from proxy
- eBPF -> { netfilter, routing }
- netfilter -> routing
- socket -> { eBPF, netfilter },



Observations & Challenges

Marking your territory

- Bitwise usage
 - Simpler interoperability
- Full-mark
 - More values to work with
 - Most usage doesn't make use of this

A tiny bit of overload


- Use every feature: 100+ bits
 - ...but there's only 32 bits to play with?
- Mitigation: Encode meaning in bit range
 - Use [0x0000..0x000F] rather than bits in 0xFFFF
- Mitigation: Overload bits on different paths
 - Ingress / Egress
 - Make semantics dependent on packet fields

Sharing is caring

- Driven by common deployment scenarios
- The clearer responsibility assignment you have, the better
- Not free (in effort or in complexity)

One does not simply understand skb mark

- Required reading: network stack diagram
- Distinct bits do not guarantee integration
 - skb, conn matches may steer packets
- Fun: replies disappear
- Proxies: Double the connections, double the fun



Less is more

- “If only I had more bits...”
- Consolidate subsystem usage
- Extend generic mark space?
- Formalize some use cases?

Summary

- Powerful mechanism for cross-subsystem programming
- Uncertainty when bits are OK to use
- There are more uses than there are bits

Cilium

- 🌐 <https://cilium.io>
- 🔗 <https://cilium.io/slack>
- 🔄 <https://github.com/cilium/cilium>
- 🐦 <https://twitter.com/ciliumproject>

Mark registry

- 🔄 <https://github.com/fwmark/registry>

